

# Deep Learning 02: *CNN*

## Lecture 11

### Computer Vision for Geosciences

2021-05-28



UNIVERSIDAD NACIONAL  
AUTÓNOMA DE  
MÉXICO

1. From MLP to CNN
2. Transfer Learning: using pretrained CNNs
3. Using TensorBoard
4. CNN cheat sheet: layers types and hyperparameters

1. From MLP to CNN
2. Transfer Learning: using pretrained CNNs
3. Using TensorBoard
4. CNN cheat sheet: layers types and hyperparameters

# Last week: MLP for MNIST-fashion dataset classification task

```
import tensorflow as tf
```

```
# Load data
fashion_mnist = tf.keras.datasets.fashion_mnist

(X_train_full, y_train_full), (X_test, y_test) = fashion_mnist.load_data()
X_valid, X_train = X_train_full[:5000], X_train_full[5000:]
y_valid, y_train = y_train_full[:5000], y_train_full[5000:]
```

```
# Preprocess data
X_train, X_test, X_valid = X_train/255.0, X_test/255.0, X_valid/255.0
```

```
# Build model (using the Sequential API)
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(300, activation='relu'),
    tf.keras.layers.Dense(100, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.summary()
```

```
# Compile model
model.compile(loss='sparse_categorical_crossentropy',
              optimizer='sgd',
              metrics=['accuracy'])
```

```
# Train model
history = model.fit(X_train, y_train, validation_data=(X_valid, y_valid),
                   epochs=30, # nb of times X_train is seen
                   batch_size=32) # nb of images per training instance
print('training instances per epoch = {}'.format(X_train.shape[0] / 32))
```

```
# Plot training history
import pandas as pd
pd.DataFrame(history.history).plot()
```

```
# Evaluate model
test_loss, test_acc = model.evaluate(X_test, y_test)
print('Test accuracy:', test_acc)
```

```
# Predict
img = X_test[0, :, :]
img = (np.expand_dims(img, 0)) # add image to a batch
y_proba = model.predict(img).round(2)
y_pred = np.argmax(model.predict(img), axis=-1)
```

```
plt.bar(range(10), y_proba[0])
plt.imshow(img[0, :, :], cmap='binary')
plt.title('class {} = {}'.format(y_pred, class_names[np.argmax(y_proba)]))
```

## 1.1 Load data

- training dataset
- validation dataset
- test dataset

## 1.2 Preprocess data

- scale pixel intensities to 0-1

## 2.1 Build model

- set layer type/order

## 2.2 Compile model

- set loss function
- set optimizer
- set metrics

## 3. Train model

- learn layer parameters (weights/biases)
- plot training history (check for overfitting)

## 4. Evaluate model

- evaluate accuracy on test dataset

## 5. Predict from model



- predict image class using learned model



Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 300)	235500
dense_1 (Dense)	(None, 100)	30100
dense_2 (Dense)	(None, 10)	1010

Total params: 266,610  
Trainable params: 266,610  
Non-trainable params: 0



# Last week: MLP for MNIST-fashion dataset classification task

```
import tensorflow as tf
```

```
# Load data
fashion_mnist = tf.keras.datasets.fashion_mnist

(X_train_full, y_train_full), (X_test, y_test) = fashion_mnist.load_data()
X_valid, X_train = X_train_full[:5000], X_train_full[5000:]
y_valid, y_train = y_train_full[:5000], y_train_full[5000:]
```

```
# Preprocess data
X_train, X_test, X_valid = X_train/255.0, X_test/255.0, X_valid/255.0
```

```
# Build model (using the Sequential API)
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(300, activation='relu'),
    tf.keras.layers.Dense(100, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.summary()
```

```
# Compile model
model.compile(loss='sparse_categorical_crossentropy',
              optimizer='sgd',
              metrics=['accuracy'])
```

```
# Train model
history = model.fit(X_train, y_train, validation_data=(X_valid, y_valid),
                   epochs=30, # nb of times X_train is seen seen
                   batch_size=32 # nb of images per training instance
                   print('training instances per epoch = {}'.format(X_train.shape[0] / 32))
```

```
# Plot training history
import pandas as pd
pd.DataFrame(history.history).plot()
```

```
# Evaluate model
test_loss, test_acc = model.evaluate(X_test, y_test)
print('Test accuracy:', test_acc)
```

```
# Predict
img = X_test[0, :, :]
img = (np.expand_dims(img, 0)) # add image to a batch
y_proba = model.predict(img).round(2)
y_pred = np.argmax(model.predict(img), axis=-1)
```

```
plt.bar(range(10), y_proba[0])
plt.imshow(img[0, :, :], cmap='binary')
plt.title('class {} = {}'.format(y_pred, class_names[np.argmax(y_proba)]))
```

## 1.1 Load data

- training dataset
- validation dataset
- test dataset



## 1.2 Preprocess data

- scale pixel intensities to 0-1

## 2.1 Build model

- set layer type/order

## 2.2 Compile model

- set loss function
- set optimizer
- set metrics

## 3. Train model



- learn layer parameters (weights/biases)
- plot training history (check for overfitting)

## 4. Evaluate model

- evaluate accuracy on test dataset

## 5. Predict from model

- predict image class using learned model



## Last week: MLP for MNIST-fashion dataset classification task

```
import tensorflow as tf
```

```
# Load data
fashion_mnist = tf.keras.datasets.fashion_mnist

(X_train_full, y_train_full), (X_test, y_test) = fashion_mnist.load_data()
X_valid, X_train = X_train_full[:5000], X_train_full[5000:]
y_valid, y_train = y_train_full[:5000], y_train_full[5000:]
```

```
# Preprocess data
X_train, X_test, X_valid = X_train/255.0, X_test/255.0, X_valid/255.0
```

```
# Build model (using the Sequential API)
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(300, activation='relu'),
    tf.keras.layers.Dense(100, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.summary()
```

```
# Compile model
model.compile(loss='sparse_categorical_crossentropy',
              optimizer='sgd',
              metrics=['accuracy'])
```

```
# Train model
history = model.fit(X_train, y_train, validation_data=(X_valid, y_valid),
                  epochs=30, # nb of times X_train is seen seen
                  batch_size=32 # nb of images per training instance
                  print('training instances per epoch = {}'.format(X_train.shape[0] / 32))
```

```
# Plot training history
import pandas as pd
pd.DataFrame(history.history).plot()
```

```
# Evaluate model
```

### 1.1 Load data

- training dataset
- validation dataset
- test dataset

### 1.2 Preprocess data

- scale pixel intensities to 0-1

### 2.1 Build model


- set layer type/order

### 2.2 Compile model

- set loss function
- set optimizer
- set metrics

### 3. Train model

- learn layer parameters (weights/biases)
- plot training history (check for overfitting)



```
plt.title('class {} - {}'.format(y_pred, class_names[np.argmax(y_proba)]))
```

MLP are powerful, but break for large images due to the huge amount of parameters to optimize

EX1: *simple* model above on the *simple* MNIST-fashion dataset (28x28 pix)  $\Rightarrow$  266,610 parameters

EX2: 100x100 image = 10,000 pixels, with first hidden layer having 1,000 neurons (which is already very restrictive)  
 $\Rightarrow 10,000 \times 1,000 = 10$  million connections, only for the first layer!

**This week:** **CNN** for MNIST-fashion dataset classification task*# Build model (MLP)*

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=[28, 28]),
    tf.keras.layers.Dense(300, activation="relu"),
    tf.keras.layers.Dense(100, activation="relu"),
    tf.keras.layers.Dense(10, activation="softmax")
])
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
flatten_5 (Flatten)	(None, 784)	0
dense_5 (Dense)	(None, 300)	235500
dense_6 (Dense)	(None, 100)	30100
dense_7 (Dense)	(None, 10)	1010
Total params: 266,610		
Trainable params: 266,610		
Non-trainable params: 0		

*# Build model (CNN)*

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(24, 7, activation="relu", padding="same", input_shape=[28, 28, 1]),
    tf.keras.layers.MaxPooling2D(2),
    tf.keras.layers.Conv2D(128, 3, activation="relu", padding="same"),
    tf.keras.layers.Conv2D(128, 3, activation="relu", padding="same"),
    tf.keras.layers.MaxPooling2D(2),
    tf.keras.layers.Conv2D(256, 3, activation="relu", padding="same"),
    tf.keras.layers.Conv2D(256, 3, activation="relu", padding="same"),
    tf.keras.layers.MaxPooling2D(2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation="relu"),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(64, activation="relu"),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(10, activation="softmax")
])
```


Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 64)	3200
max_pooling2d (MaxPooling2D)	(None, 14, 14, 64)	0
conv2d_1 (Conv2D)	(None, 14, 14, 128)	73856
conv2d_2 (Conv2D)	(None, 14, 14, 128)	147584
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 128)	0
conv2d_3 (Conv2D)	(None, 7, 7, 256)	295168
conv2d_4 (Conv2D)	(None, 7, 7, 256)	590880
max_pooling2d_2 (MaxPooling2D)	(None, 3, 3, 256)	0
flatten (Flatten)	(None, 2304)	0
dense (Dense)	(None, 128)	295040
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 64)	8256
dropout_1 (Dropout)	(None, 64)	0
dense_2 (Dense)	(None, 10)	650
Total params: 1,413,834		
Trainable params: 1,413,834		
Non-trainable params: 0		

**This week:** CNN for MNIST-fashion dataset classification task


```
# Build model (MLP)
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=[28, 28]),
    tf.keras.layers.Dense(300, activation="relu"),
    tf.keras.layers.Dense(100, activation="relu"),
    tf.keras.layers.Dense(10, activation="softmax")
])
```

Multi Layer Perceptron (MLP)



```
# Build model (CNN)
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(64, 7, activation="relu", padding="same", input_shape=[28, 28, 1]),
    tf.keras.layers.MaxPooling2D(2),
    tf.keras.layers.Conv2D(128, 3, activation="relu", padding="same"),
    tf.keras.layers.Conv2D(128, 3, activation="relu", padding="same"),
    tf.keras.layers.MaxPooling2D(2),
    tf.keras.layers.Conv2D(256, 3, activation="relu", padding="same"),
    tf.keras.layers.Conv2D(256, 3, activation="relu", padding="same"),
    tf.keras.layers.MaxPooling2D(2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation="relu"),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(64, activation="relu"),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(10, activation="softmax")
])
```

Convolutional Neural Network (CNN)





1. From MLP to CNN
2. Transfer Learning: using pretrained CNNs
3. Using TensorBoard
4. CNN cheat sheet: layers types and hyperparameters

Designing and training your own network can be difficult (or impossible if you do not have enough data)

⇒ Transfer Learning allows to fine-tune a pretrained network

⇒ Most famous CNN networks achieving very good performances on the [ImageNet](#) dataset:

*(ImageNet = several millions of images, large size (256 pixels), with >1000 classes)*

- LeNet-5 (1998)
- AlexNet (2012)
- GoogLeNet (2014)
- ResNet (2015)
- Xception (2016)
- SENet (2017)

Designing and training your own network can be difficult (or impossible if you do not have enough data)

⇒ Transfer Learning allows to fine-tune a pretrained network

⇒ Most famous CNN networks achieving very good performances on the [ImageNet](#) dataset:

*(ImageNet = several millions of images, large size (256 pixels), with > 1000 classes)*

- LeNet-5 (1998)
- AlexNet (2012)
- GoogLeNet (2014)
- ResNet (2015)
- Xception (2016)
- SNet (2017)

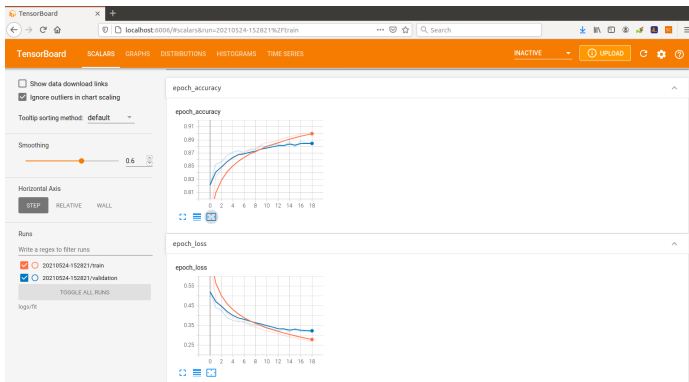
In the today's exercise we will use the **Xception** model with weights learned from the **ImageNet** dataset

1. From MLP to CNN
2. Transfer Learning: using pretrained CNNs
- 3. Using TensorBoard**
4. CNN cheat sheet: layers types and hyperparameters

## TensorBoard: TensorFlow's visualization toolkit

⇒ TensorBoard provides the visualization and tooling needed for machine learning experimentation:

- Tracking and visualizing metrics such as loss and accuracy
- Visualizing the model graph (ops and layers)
- Viewing histograms of weights, biases
- etc.



### TensorBoard: TensorFlow's visualization toolkit

⇒ TensorBoard is installed during the TensorFlow conda installation

⇒ To use it, you should:

1. Add the `tf.keras.callbacks.TensorBoard` callback to the Keras `Model.fit()` method (ensures that logs are created and stored)

```
# Create callback
import datetime
log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1)

# Add callback to model.fit()
history = model.fit(X_train, y_train, callbacks=[tensorboard_callback])
```

2. Run TensorBoard from command line

```
# conda activate tf
# cd <working dir>
# tensorboard --logdir logs/fit # set directory used to store logs
```

3. Open a web-browser to the address

```
http://localhost:6006/
```

Nota Bene: you can open it directly from a Jupyter cell (after training has finished however) as follows:

```
%load_ext tensorboard # Load the TensorBoard notebook extension
%tensorboard --logdir logs # Open TensorBoard in cell
```

**TensorBoard:** TensorFlow's visualization toolkit

⇒ TensorBoard is installed during the TensorFlow conda installation

⇒ To use it, you should:

1. Add the **tf.keras.callbacks.TensorBoard** callback to the Keras Model.fit() method *(ensures that logs are created and stored)*

```
# Create callback
import datetime
log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1)

# Add callback to model.fit()
history = model.fit(X_train, y_train, callbacks=[tensorboard_callback])
```

2. Run TensorBoard from command line

```
$ conda activate tf
$ cd <working dir>
$ tensorboard --logdir logs/fit # set directory used to store logs
```

3. Open a web-browser to the address

```
http://localhost:6006/
```

**Nota Bene:** you can open it directly from a Jupyter cell (after training has finished however) as follows:

```
%load_ext tensorboard # Load the TensorBoard notebook extension
%tensorboard --logdir logs # Open TensorBoard in cell
```

## TensorBoard: TensorFlow's visualization toolkit

⇒ TensorBoard is installed during the TensorFlow conda installation

⇒ To use it, you should:

1. Add the **tf.keras.callbacks.TensorBoard** callback to the Keras Model.fit() method *(ensures that logs are created and stored)*

```
# Create callback
import datetime
log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1)

# Add callback to model.fit()
history = model.fit(X_train, y_train, callbacks=[tensorboard_callback])
```

2. Run TensorBoard from command line

```
$ conda activate tf
$ cd <working dir>
$ tensorboard --logdir logs/fit # set directory used to store logs
```

3. Open a web-browser to the address

```
http://localhost:6006/
```

Nota Bene: you can open it directly from a Jupyter cell (after training has finished however) as follows:

```
%load_ext tensorboard # Load the TensorBoard notebook extension
%tensorboard --logdir logs # Open TensorBoard in cell
```



## TensorBoard: TensorFlow's visualization toolkit

⇒ TensorBoard is installed during the TensorFlow conda installation

⇒ To use it, you should:

1. Add the **tf.keras.callbacks.TensorBoard** callback to the Keras Model.fit() method *(ensures that logs are created and stored)*

```
# Create callback
import datetime
log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1)

# Add callback to model.fit()
history = model.fit(X_train, y_train, callbacks=[tensorboard_callback])
```

2. Run TensorBoard from command line

```
$ conda activate tf
$ cd <working dir>
$ tensorboard --logdir logs/fit # set directory used to store logs
```

3. Open a web-browser to the address

```
http://localhost:6006/
```

Nota Bene: you can open it directly from a Jupyter cell (after training has finished however) as follows:

```
%load_ext tensorboard # Load the TensorBoard notebook extension
%tensorboard --logdir logs # Open TensorBoard in cell
```

## TensorBoard: TensorFlow's visualization toolkit

⇒ TensorBoard is installed during the TensorFlow conda installation

⇒ To use it, you should:

1. Add the **tf.keras.callbacks.TensorBoard** callback to the Keras Model.fit() method (*ensures that logs are created and stored*)

```
# Create callback
import datetime
log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1)

# Add callback to model.fit()
history = model.fit(X_train, y_train, callbacks=[tensorboard_callback])
```

2. Run TensorBoard from command line

```
$ conda activate tf
$ cd <working dir>
$ tensorboard --logdir logs/fit # set directory used to store logs
```

3. Open a web-browser to the address

```
http://localhost:6006/
```

Nota Bene: you can open it directly from a Jupyter cell (after training has finished however) as follows:


```
%load_ext tensorboard # Load the TensorBoard notebook extension
%tensorboard --logdir logs # Open TensorBoard in cell
```

1. From MLP to CNN
2. Transfer Learning: using pretrained CNNs
3. Using TensorBoard
4. CNN cheat sheet: layers types and hyperparameters

# CNN cheat sheet: layers types and hyperparameters

## CONV = convolutional layer

=> 1 filter = "boxcar" filter (where each pixel is multiplied by a weight, products summed across bands, and bias added to the result), with sliding step S  
=> convolution layer = N filters




EXAMPLE: Input volume:  
size  $h_{in}=5, d_{in}=3$ , zero-padding  $P=0$   
=> X.shape =  $(h_{in}+2P, h_{in}+2P, d_{in}) = (5,5,3)$

Filter kernel:  
depth  $K=2$  ( $W0$  &  $W1$ ), size  $F=3$ , stride  $S=1$   
=>  $W0$ .shape =  $(F, F, d_{in}) = (3,3,3)$   
=>  $W0$  weights =  $F^2 d_{in} = 3^2 = 27$ ,  $W0$  bias = 1

Output volume:  
 $h_{out} = (h_{in} - F + 2P) / S + 1 = (5 - 3 + 0) / 1 + 1 = 3$   
=> V.shape =  $(3,3,2)$


## ReLU = activation layer

=> activation function  
=> applies threshold



## POOL = max pool layer

=> reduces dimensionality (for memory issue)  
=> enables network to see image at large scale



## FC = fully connected

## Hyper parameters

### - receptive field (F)

= filter size

NB: usually an odd number, so that it is centered on a central pixel

### - depth (K)

= number of filters

NB: depth column = set of neurons that are all looking at the same region of the input

### - stride (S)

= number of pixels the filter slides across the image at each step

EX: stride 2 => filter moves 2 pix at a time => produces smaller outputs

### - zero-padding (P)

= pad the input volume with zeros around the border

NB: hyper-parameters control the output volume size:

width & height =  $(W - F + 2P) / S + 1$  where  $W$  = input width/height  
depth =  $K$