

Digital Image Basics

Lecture 02

Computer Vision for Geosciences

2021-03-05




UNIVERSIDAD NACIONAL
AUTÓNOMA DE
MÉXICO

1. Motivation
 - sources of images
2. What is a digital image?
 - eye versus pinhole camera
 - sampling and quantization
 - color image
 - color spaces
 - image histogram
3. Point operations
 - homogeneous point operations
 - inhomogeneous Point Operations
4. Computer Vision
 - categorizing processing tasks
5. Image manipulation with Python
 - numpy tutorial + exercises


1. Motivation
 - sources of images
2. What is a digital image?
 - eye versus pinhole camera
 - sampling and quantization
 - color image
 - color spaces
 - image histogram
3. Point operations
 - homogeneous point operations
 - inhomogeneous Point Operations
4. Computer Vision
 - categorizing processing tasks
5. Image manipulation with Python
 - numpy tutorial + exercises

Images can be constructed using the entire electromagnetic spectra




... a few examples in geosciences ...

Images can be constructed using the entire electromagnetic spectra



... a few examples in geosciences ...

Motivation: sources of images




camera




Popocatepetl 2020-04-16

satellite




Popocatepetl 2021-02-25 (Sentinel-2, MOUNTS)

microscope




Popocatepetl 2019-01-22 (andesite, ©T.Boulesteix)

Motivation: sources of images




camera




Popocatepetl 2013-01-29 (UV camera, [Campion et al. 2018](#))

satellite




Popocatepetl 2019-02-17 (Sentinel-5P, [MOUNTS](#))

Motivation: sources of images




camera




Nyiragongo 2016-04-16 (FLIR image, [Valade et al. 2018](#))

satellite




Etna 2021-02-23 (Sentinel-2 image, [MOUNTS](#))

Motivation: sources of images




satellite (SAR)




Popocatepetl 2021-02-28 (Sentinel-1, [MOUNTS](#))

satellite (InSAR)




Popocatepetl InSAR interferogram ([MOUNTS](#))

satellite (InSAR)





Popocatepetl InSAR coherence ([MOUNTS](#)) 9 / 61

Motivation: sources of images




telescope




Crab Nebula - remanent of an exploded star (supernova)

Motivation: sources of images



telescope




Radio wave (VLA)




Infrared radiation (Spitzer)



Visible light (Hubble)



Ultraviolet radiation (Astro-1)




Low-energy X-ray (Chandra)




High-energy X-ray (HEFT)
*** 15 min exposure ***

Crab Nebula - remnant of an exploded star (supernova)

Motivation: sources of images



telescope




Radio wave (VLA)



Infrared radiation (Spitzer)



Visible light (Hubble)



Ultraviolet radiation (Astro-1)



Low-energy X-ray (Chandra)



Pixel Size

High-energy X-ray (HEFT)
*** 15 min exposure ***

Crab Nebula - remanent of an exploded star (supernova)

1. Motivation
 - sources of images
2. What is a digital image?
 - eye versus pinhole camera
 - sampling and quantization
 - color image
 - color spaces
 - image histogram
3. Point operations
 - homogeneous point operations
 - inhomogeneous Point Operations
4. Computer Vision
 - categorizing processing tasks
5. Image manipulation with Python
 - numpy tutorial + exercises

Comparison between human eye and pinhole camera





Image = 3D world projection on 2D

⇒ projection using the **pinhole camera** model:



Perspective transformation:

$$s m' = K[R|t]M' \quad (1)$$


$$s \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & 0 & u_0 \\ 0 & f_y & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_1 \\ r_{21} & r_{22} & r_{23} & t_2 \\ r_{31} & r_{32} & r_{33} & t_3 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \quad (2)$$

where:

- M' = 3D point in space with coordinates $[X, Y, Z]^T$ expressed in Euclidean coordinates
- m' = projection of the 3D point M' onto the image plane with coordinates $[u, v]^T$ expressed in pixel units
- K = camera calibration matrix (a.k.a. intrinsic parameters matrix)
 - f_x, f_y = focal lengths expressed in pixel units
 - u_0, v_0 = coordinates of the optical center (aka principal point), origin in the image plane
- $[R|t]$ = joint rotation-translation matrix (a.k.a. extrinsic parameters matrix), describing the camera pose, and translating from world coordinates to camera coordinates

Image = 3D world projection on 2D

⇒ projection using the **pinhole camera** model:



Perspective transformation:

$$s m' = K[R|t]M' \quad (1)$$

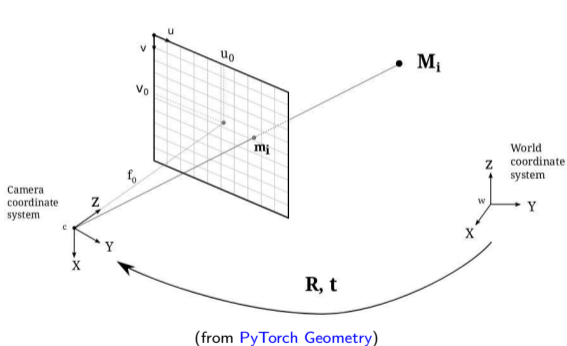
$$s \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & 0 & u_0 \\ 0 & f_y & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_1 \\ r_{21} & r_{22} & r_{23} & t_2 \\ r_{31} & r_{32} & r_{33} & t_3 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \quad (2)$$

where:

- M' = 3D point in space with coordinates $[X, Y, Z]^T$ expressed in Euclidean coordinates
- m' = projection of the 3D point M' onto the image plane with coordinates $[u, v]^T$ expressed in pixel units
- K = camera calibration matrix (a.k.a. intrinsics parameters matrix)
 - f_x, f_y = focal lengths expressed in pixel units
 - u_0, v_0 = coordinates of the optical center (aka principal point), origin in the image plane
- $[R|t]$ = joint rotation-translation matrix (a.k.a. extrinsics parameters matrix), describing the camera pose, and translating from world coordinates to camera coordinates

Image = 3D world projection on 2D

⇒ projection using the **pinhole camera** model:



Perspective transformation:

$$s m' = K[R|t]M' \quad (1)$$


$$s \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & 0 & u_0 \\ 0 & f_y & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_1 \\ r_{21} & r_{22} & r_{23} & t_2 \\ r_{31} & r_{32} & r_{33} & t_3 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \quad (2)$$

where:

- M' = 3D point in space with coordinates $[X, Y, Z]^T$ expressed in Euclidean coordinates
- m' = projection of the 3D point M' onto the image plane with coordinates $[u, v]^T$ expressed in pixel units
- K = camera calibration matrix (a.k.a. intrinsics parameters matrix)
 - f_x, f_y = focal lengths expressed in pixel units
 - u_0, v_0 = coordinates of the optical center (aka principal point), origin in the image plane
- $[R|t]$ = joint rotation-translation matrix (a.k.a. extrinsics parameters matrix), describing the camera pose, and translating from world coordinates to camera coordinates

Image = 3D world projection on 2D

⇒ projection using the **pinhole camera** model:



Perspective transformation:

$$s m' = K[R|t]M' \quad (1)$$


$$s \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & 0 & u_0 \\ 0 & f_y & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_1 \\ r_{21} & r_{22} & r_{23} & t_2 \\ r_{31} & r_{32} & r_{33} & t_3 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \quad (2)$$

where:

- M' = 3D point in space with coordinates $[X, Y, Z]^T$ expressed in Euclidean coordinates
- m' = projection of the 3D point M' onto the image plane with coordinates $[u, v]^T$ expressed in pixel units
- K = camera calibration matrix (a.k.a intrinsics parameters matrix)
 - f_x, f_y = focal lengths expressed in pixel units
 - u_0, v_0 = coordinates of the optical center (aka principal point), origin in the image plane
- $[R|t]$ = joint rotation-translation matrix (a.k.a. extrinsics parameters matrix), describing the camera pose, and translating from world coordinates to camera coordinates

Image = 3D world projection on 2D

⇒ projection using the **pinhole camera** model:



Perspective transformation:


$$s m' = K[R|t]M' \quad (1)$$

$$s \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & 0 & u_0 \\ 0 & f_y & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_1 \\ r_{21} & r_{22} & r_{23} & t_2 \\ r_{31} & r_{32} & r_{33} & t_3 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \quad (2)$$


where:

- M' = 3D point in space with coordinates $[X, Y, Z]^T$ expressed in Euclidean coordinates
- m' = projection of the 3D point M' onto the image plane with coordinates $[u, v]^T$ expressed in pixel units
- K = camera calibration matrix (a.k.a intrinsics parameters matrix)
 - f_x, f_y = focal lengths expressed in pixel units
 - u_0, v_0 = coordinates of the optical center (aka principal point), origin in the image plane
- $[R|t]$ = joint rotation-translation matrix (a.k.a. extrinsics parameters matrix), describing the camera pose, and translating from world coordinates to camera coordinates

- at each point we record incident light
- digitalization of an analog signal involves two operations
 - spatial sampling (= discretization of space domain)
 - intensity quantization (= discretization of incoming light signal)



- at each point we record incident light
- digitalization of an analog signal involves two operations
 - **spatial sampling** (= discretization of space domain)
 - **intensity quantization** (= discretization of incoming light signal)



spatial sampling (= discretization of space domain)

⇒ smallest element resulting from the discretization of the space is called a pixel (=picture element)

(512, 512)



(128, 128)



(64, 64)



(32, 32)

**intensity quantization** (= discretization of light intensity signal)

⇒ typically, 256 levels (8 bits/pixel = 2^8 values) suffices to represent the intensity

8-bit resolution
 $2^8 = 256$ gray levels



3-bit resolution
 $2^3 = 8$ gray levels



2-bit resolution
 $2^2 = 4$ gray levels



1-bit resolution
 $2^1 = 2$ gray levels



spatial sampling (= discretization of space domain)

⇒ smallest element resulting from the discretization of the space is called a pixel (=picture element)

(512, 512)



(128, 128)



(64, 64)



(32, 32)

**intensity quantization** (= discretization of light intensity signal)

⇒ typically, 256 levels (8 bits/pixel = 2^8 values) suffices to represent the intensity

8-bit resolution
 $2^8 = 256$ gray levels



3-bit resolution
 $2^3 = 8$ gray levels




2-bit resolution
 $2^2 = 4$ gray levels



1-bit resolution
 $2^1 = 2$ gray levels



⇒ digital image function $f(x, y)$



⇒ digital image function $f(x, y)$



	columns									
	0	1	2	3	4	5	6	7	8	9
0	12	34	67	103	87	79	176	138	94	180
1	98	53	66	226	11	34	241	240	24	143
2	228	107	60	58	144	251	137	93	86	130
3	155	108	132	159	129	141	245	211	100	1
4	91	187	67	135	49	175	193	61	24	183
5	199	251	80	1	121	105	222	147	226	63
6	181	27	56	238	113	158	176	47	167	109
7	36	172	18	192	184	162	181	202	17	72
8	13	106	30	11	53	68	178	232	91	219
9	211	181	78	1	11	185	204	106	131	70

Typical ranges:

- `uint8` = [0-255]
(8 bits = 1 byte = $2^8 = 256$ values per pixel)
- `float32` = [0-1]
(32 bits = 4 bytes = $4.3e9$ values per pixel)

How do we record colors?

⇒ **Bayer Filter**: color filter array for arranging RGB color filters on a square grid of photosensors



(source [wikipedia](#))

How do we record colors?

⇒ **Bayer Filter**: color filter array for arranging RGB color filters on a square grid of photosensors



1. Original scene
2. Output of a 120×80 -pixel sensor with a Bayer filter
3. Output color-coded with Bayer filter colors
4. Reconstructed image after interpolating missing color information (a.k.a. demosaicing)
5. Full RGB version at 120×80 -pixels for comparison

How do we record colors?


⇒ **Bayer Filter**: color filter array for arranging RGB color filters on a square grid of photosensors



1. Original scene
2. Output of a 120×80 -pixel sensor with a Bayer filter
3. Output color-coded with Bayer filter colors
4. Reconstructed image after interpolating missing color information (a.k.a. demosaicing)
5. Full RGB version at 120×80 -pixels for comparison

How do we record colors?


⇒ **Bayer Filter**: color filter array for arranging RGB color filters on a square grid of photosensors



1. Original scene
2. Output of a 120×80 -pixel sensor with a Bayer filter
3. Output color-coded with Bayer filter colors
4. Reconstructed image after interpolating missing color information (a.k.a. demosaicing)
5. Full RGB version at 120×80 -pixels for comparison

How do we record colors?


⇒ **Bayer Filter**: color filter array for arranging RGB color filters on a square grid of photosensors



1. Original scene
2. Output of a 120×80 -pixel sensor with a Bayer filter
3. Output color-coded with Bayer filter colors
4. Reconstructed image after interpolating missing color information (a.k.a. demosaicing)
5. Full RGB version at 120×80 -pixels for comparison

How do we record colors?

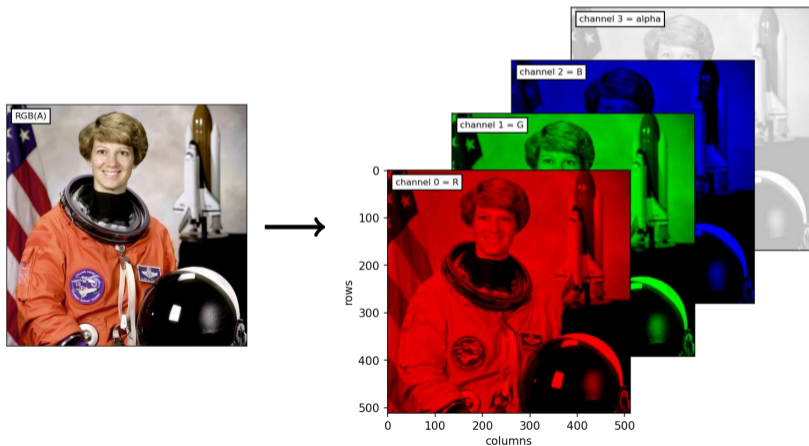
⇒ **Bayer Filter**: color filter array for arranging RGB color filters on a square grid of photosensors



1. Original scene
2. Output of a 120×80 -pixel sensor with a Bayer filter
3. Output color-coded with Bayer filter colors
4. Reconstructed image after interpolating missing color information (a.k.a. demosaicing)
5. Full RGB version at 120×80 -pixels for comparison


⇒ color image = 3D tensor in colorspace

- **RGB** = Red + Green + Blue bands (.JPEG)
- **RGBA** = Red + Green + Blue + Alpha bands (.PNG, .GIF, .BMP, TIFF, .JPEG 2000)




Other ways to represent the color information?

RGB colorspace




HSV colorspace



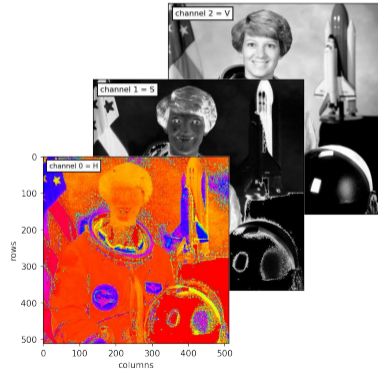
- Hue (H) = $[0-360]$ \Rightarrow shift color
- Saturation (S) = $[0-1]$ \Rightarrow shift intensity
- Value (V) = $[0-1]$ \Rightarrow shift brightness

3D tensor with different information

RGB colorspace



HSV colorspace



- more saturation S

⇒ more intense colors



- more value V

⇒ brighter colors

- shift hue H

⇒ shift color

- more saturation S

⇒ more intense colors

original



saturation x2



- more value V

⇒ brighter colors

original



value x1.5



- shift hue H

⇒ shift color

- more saturation S

⇒ more intense colors

original



saturation x2



original



value x1.5



- more value V

⇒ brighter colors

original



hue x5




- shift hue H


⇒ shift color

Histogram of pixel values in each band:

original (uint8)




Histogram of pixel values after conversion from RGB (3-bands) to gray-scale (1-band):




Histogram of pixel values after conversion to float values (range [0-1])

gray-scale (float)




- original gray-scale




- histogram rescale to 10-90 percentiles
⇒ contrast stretching

- histogram equalize
⇒ spread out the most frequent intensity values

- original gray-scale




- histogram rescale to 10-90 percentiles
⇒ contrast stretching




- histogram equalize
⇒ spread out the most frequent intensity values


- original gray-scale



- histogram rescale to 10-90 percentiles
⇒ contrast stretching

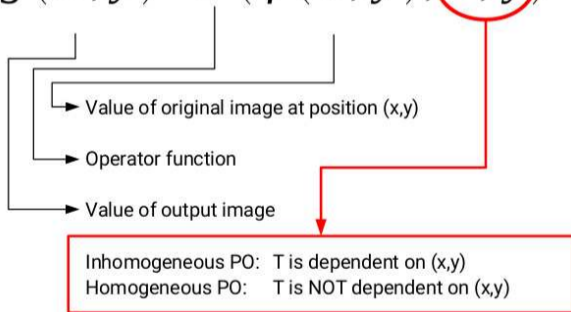


- histogram equalize
⇒ spread out the most frequent intensity values



1. Motivation
 - sources of images
2. What is a digital image?
 - eye versus pinhole camera
 - sampling and quantization
 - color image
 - color spaces
 - image histogram
3. Point operations
 - homogeneous point operations
 - inhomogeneous Point Operations
4. Computer Vision
 - categorizing processing tasks
5. Image manipulation with Python
 - numpy tutorial + exercises

$$g(x, y) = T(f(x, y), x, y)$$



Point operations


1. homogeneous point operations

Homogeneous Point Operations (does not depend on pixel position)

identity



$g(x, y) = f(x, y)$



Point operations

1. homogeneous point operations

Homogeneous Point Operations (does not depend on pixel position)


identity



inverse



$$g(x, y) = 1 - f(x, y)$$



Homogeneous Point Operations (does not depend on pixel position)

identity




inverse



square root



$$g(x, y) = \sqrt{f(x, y)}$$



Point operations

1. homogeneous point operations

Homogeneous Point Operations (does not depend on pixel position)

identity



inverse




square root



logarithm



$$g(x, y) = a \cdot \log(f(x, y) + 1)$$



Point operations

1. homogeneous point operations

Homogeneous Point Operations (does not depend on pixel position)

identity



inverse



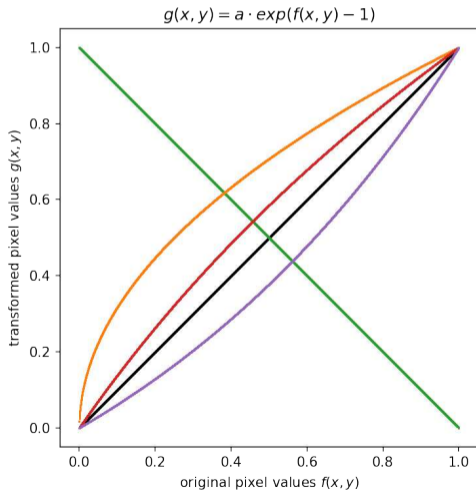
square root



logarithm



exponential



Homogeneous Point Operations (does not depend on pixel position)

identity



inverse



square root




logarithm



exponential



square



Inhomogeneous Point Operations (depends on pixel position)

EX: background detection / change detection

 f_1  f_i  f_N

$$a(x, y) = \frac{1}{N} \sum_{i=0}^N f_i(x, y)$$


 $a(x, y)$

$$\begin{aligned} g_i(x, y) &= T(f(x, y), x, y) \\ &= f_i(x, y) - a(x, y) \end{aligned}$$



Inhomogeneous Point Operations (depends on pixel position)

EX: background detection / change detection



1. Motivation
 - sources of images
2. What is a digital image?
 - eye versus pinhole camera
 - sampling and quantization
 - color image
 - color spaces
 - image histogram
3. Point operations
 - homogeneous point operations
 - inhomogeneous Point Operations
4. **Computer Vision**
 - categorizing processing tasks
5. Image manipulation with Python
 - numpy tutorial + exercises

Computer Vision processing levels:

- Low-level vision
 - image manipulation
(resizing, color adjustments, ...)
 - feature extraction
(edges, gradients, ...)
- Mid-level vision
 - panorama stitching
 - Structure from Motion (SfM) \Rightarrow 2D to 3D
 - Optical Flow \Rightarrow velocities
- High-level vision
 - classification: what is in the image?
 - tagging: what are ALL the things in the image?
 - detection: where are they?
 - semantic segmentation \Rightarrow segment image and give names

Computer Vision processing levels:

- Low-level vision
 - image manipulation
(resizing, color adjustments, ...)
 - feature extraction
(edges, gradients, ...)
- Mid-level vision
 - panorama stitching
 - Structure from Motion (SfM) \Rightarrow 2D to 3D
 - Optical Flow \Rightarrow velocities
- High-level vision
 - classification: what is in the image?
 - tagging: what are ALL the things in the image?
 - detection: where are they?
 - semantic segmentation \Rightarrow segment image and give names

Computer Vision processing levels:

- Low-level vision
 - image manipulation
(resizing, color adjustments, ...)
 - feature extraction
(edges, gradients, ...)
- Mid-level vision
 - panorama stitching
 - Structure from Motion (SfM) \Rightarrow 2D to 3D
 - Optical Flow \Rightarrow velocities
- High-level vision
 - classification: what is in the image?
 - tagging: what are ALL the things in the image?
 - detection: where are they?
 - semantic segmentation \Rightarrow segment image and give names

Computer Vision processing levels:

- Low-level vision
 - image manipulation
(resizing, color adjustments, ...)
 - feature extraction
(edges, gradients, ...)
- Mid-level vision
 - panorama stitching
 - Structure from Motion (SfM) \Rightarrow 2D to 3D
 - Optical Flow \Rightarrow velocities
- High-level vision
 - classification: what is in the image?
 - tagging: what are ALL the things in the image?
 - detection: where are they?
 - semantic segmentation \Rightarrow segment image and give names

1. Motivation
 - sources of images
2. What is a digital image?
 - eye versus pinhole camera
 - sampling and quantization
 - color image
 - color spaces
 - image histogram
3. Point operations
 - homogeneous point operations
 - inhomogeneous Point Operations
4. Computer Vision
 - categorizing processing tasks
5. Image manipulation with Python
 - numpy tutorial + exercises

In Binder:

⇒ Open CV4GS_02_imagebasics/[CV4GS_02_numpy-tutorial.ipynb](#)

⇒ Open CV4GS_02_imagebasics/[CV4GS_02_exercices.ipynb](#)

In Binder:

⇒ Open CV4GS_02_imagebasics/[CV4GS_02_numpy-tutorial.ipynb](#)

⇒ Open CV4GS_02_imagebasics/[CV4GS_02_exercices.ipynb](#)