

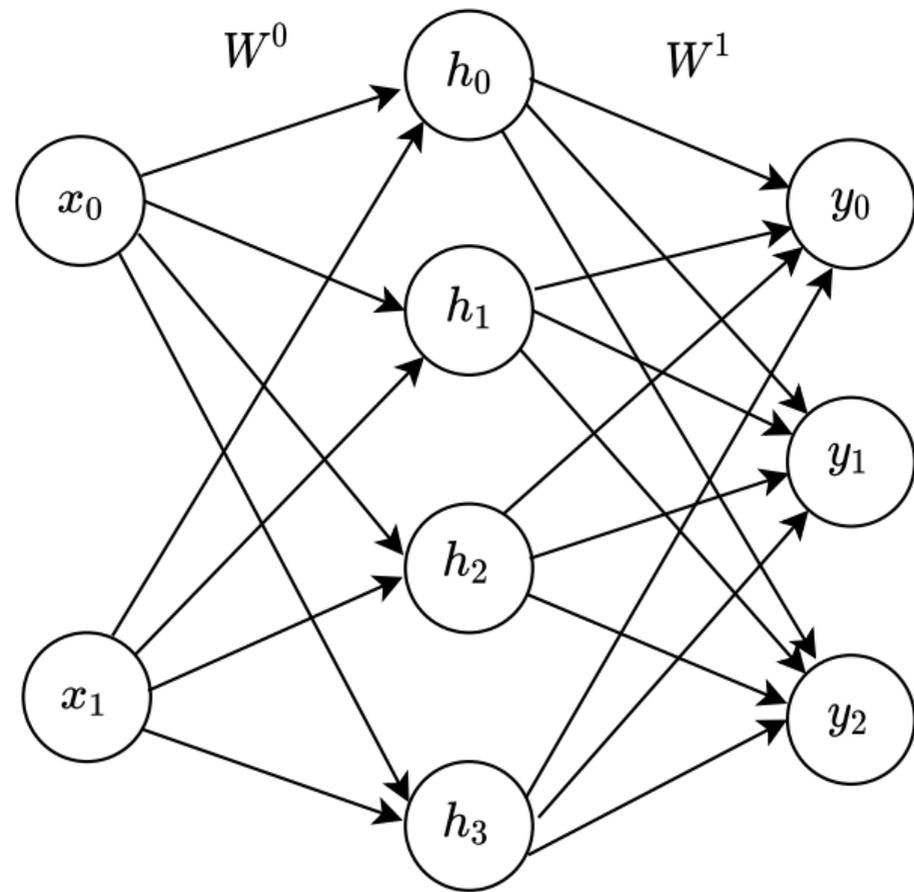
Neural Networks 2/3

Lecture 11

Computer Vision for Geosciences

May 28, 2021





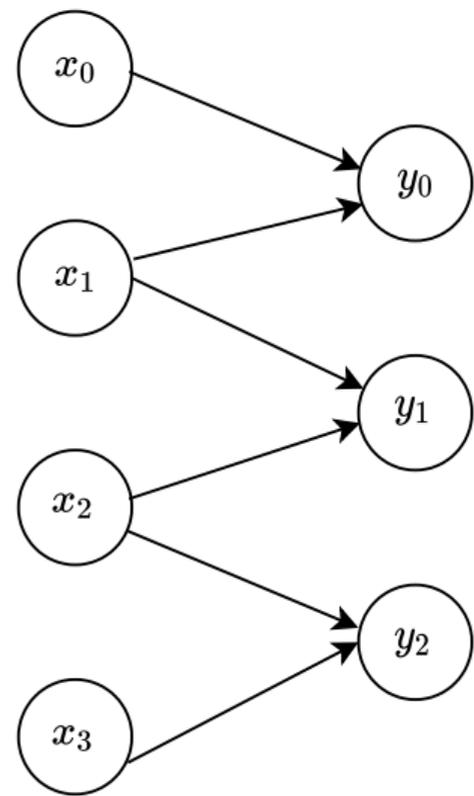
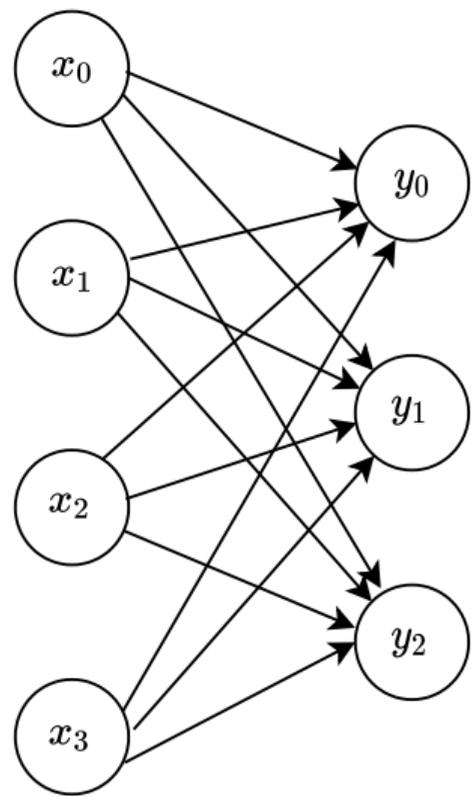
- Input: an rgb image relatively low resolution
→ $128 \times 128 \times 3 = 49152$
- Hidden layer: comparable to HOG with 36 dim feature vector computed from 8x8 patches
→ $16 \times 16 \times 36 = 9216$
- Output neurons for e.g. 10 object classes
→ $49152 \cdot 9216 + 9216 \cdot 10 = 453076992 \approx 450$ million parameters



Do we need to connect all the pixels?

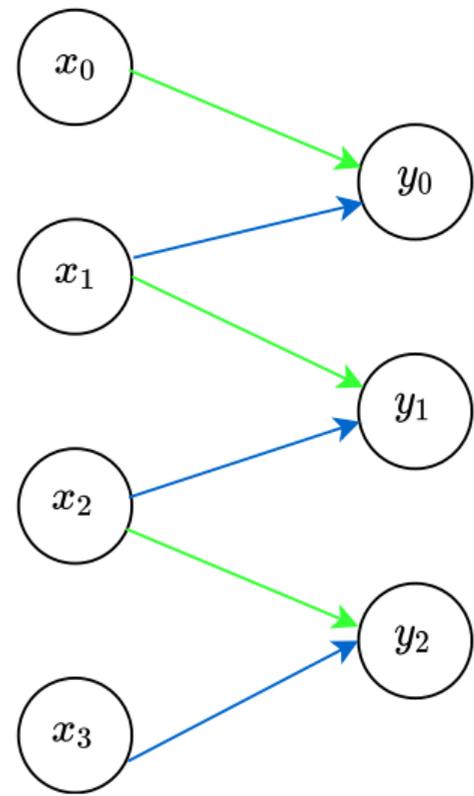
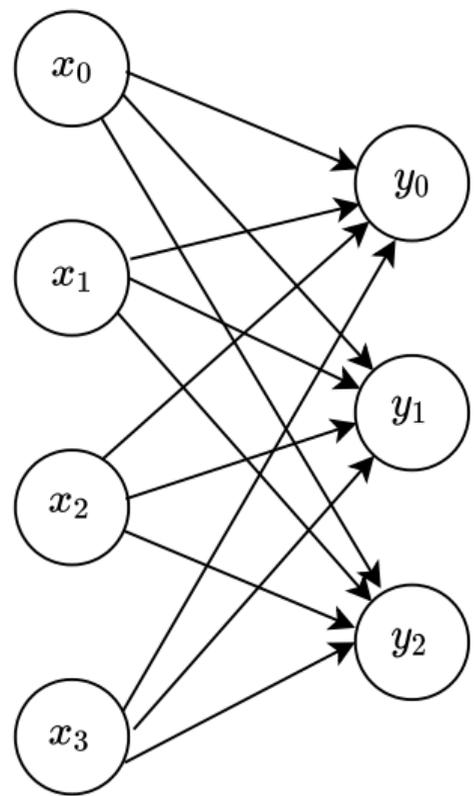
- Too many weights.
- Local image statistics.
- Weight sharing.

Convolutional Layers: locality



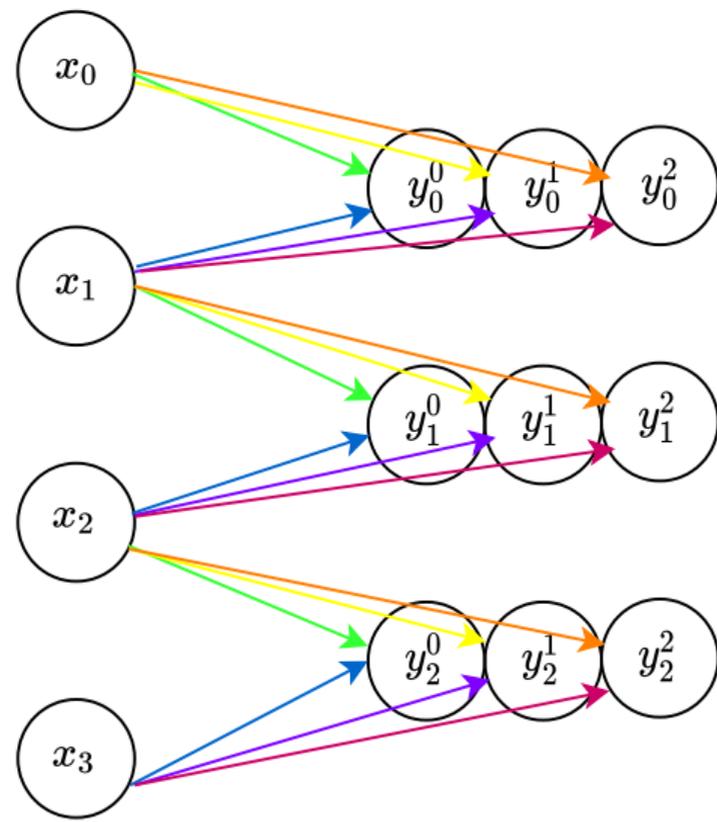
- Assumption: local regions to be processed together, regions far apart not related

Convolutional Layers: weight sharing



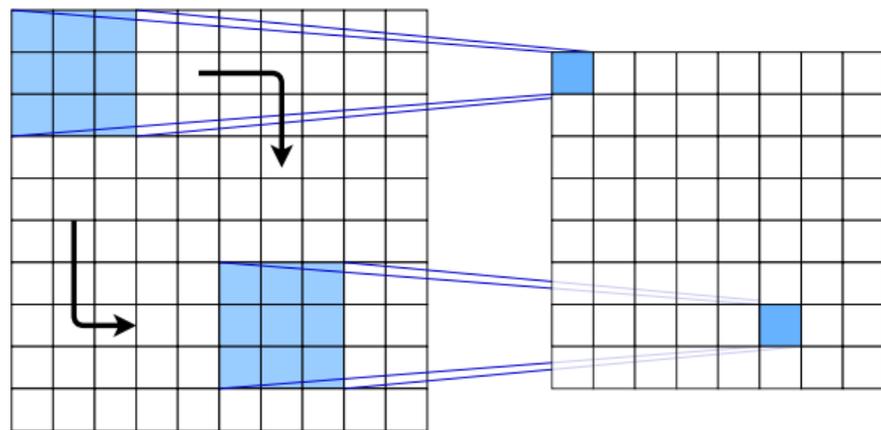
- Assumption: image processing should not vary with image region.

Convolutional Layers: feature maps

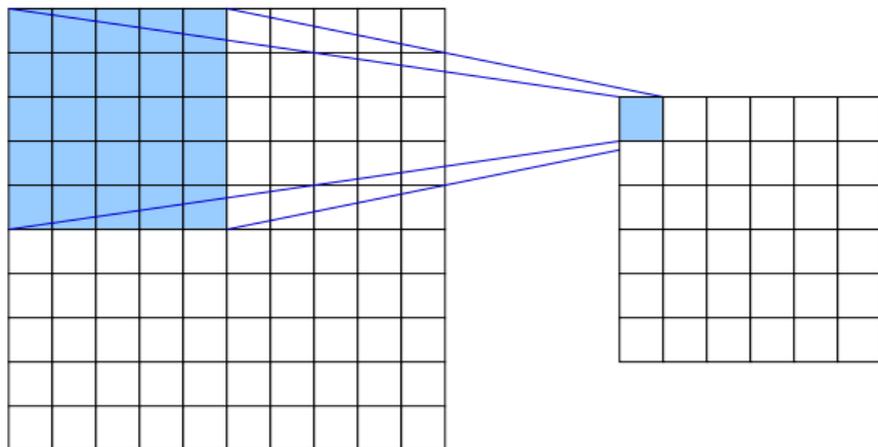


- Instead we can connect multiple neurons to every dimension of the input.

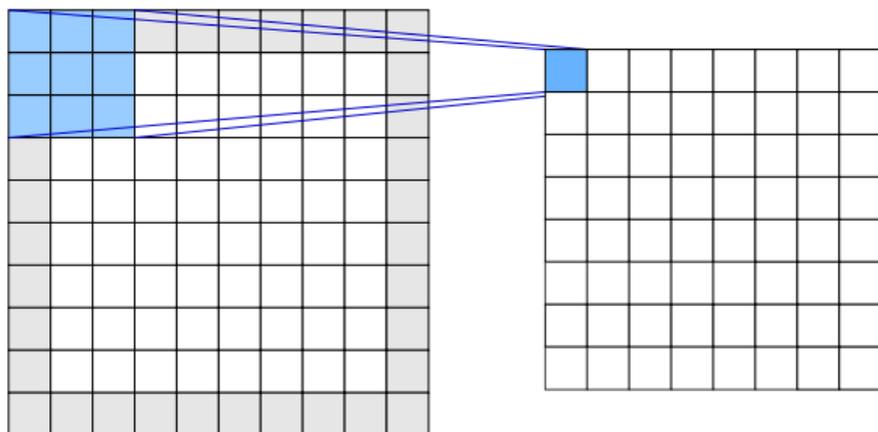
- The perceptron operation $\sigma(Wx)$ corresponds to a convolution with a filter kernel plus non linearity.



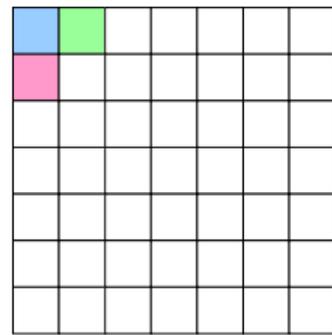
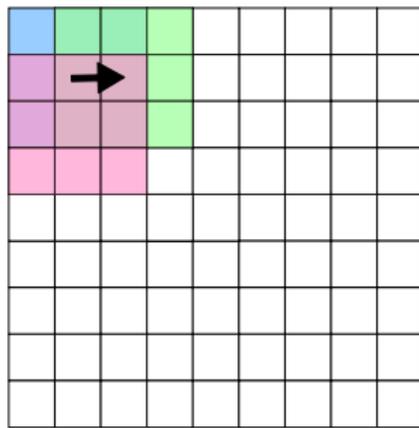
- We lose $\frac{1}{2}$ kernel size pixels at the image boarder.



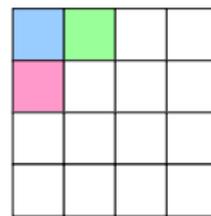
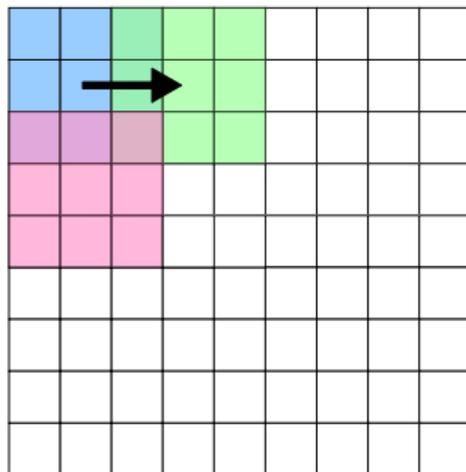
- Usually we pad image boarder to keep image size.



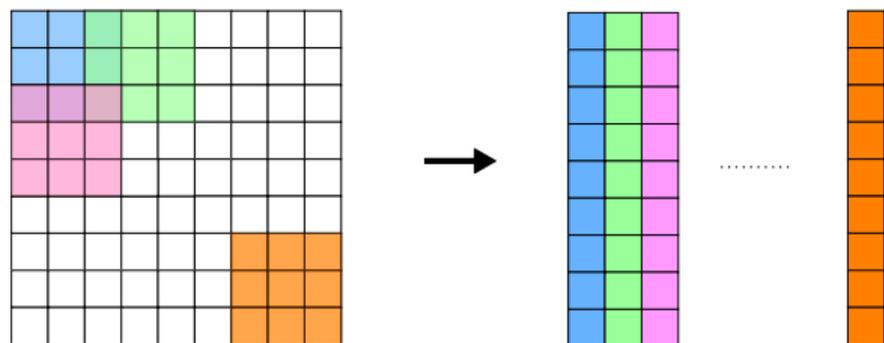
- The number of pixels in between neurons is called stride.

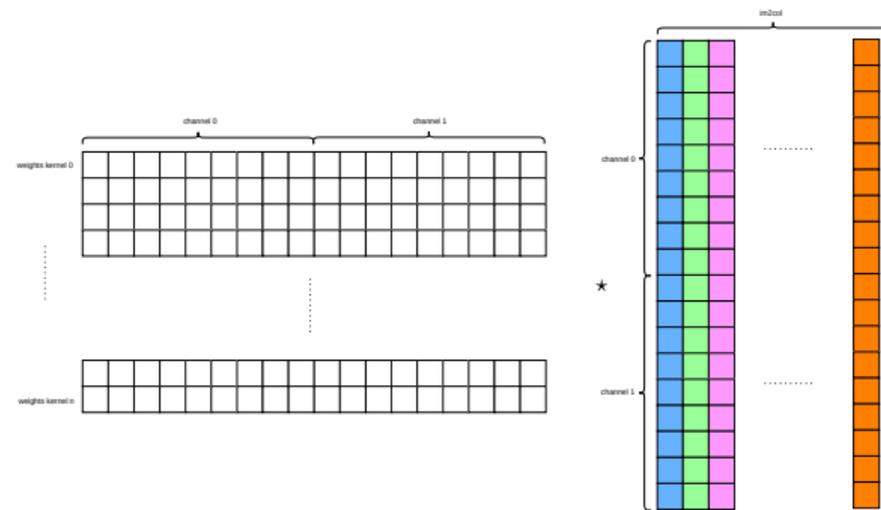


- With strides > 1 we can downsample the image to lower resolution.



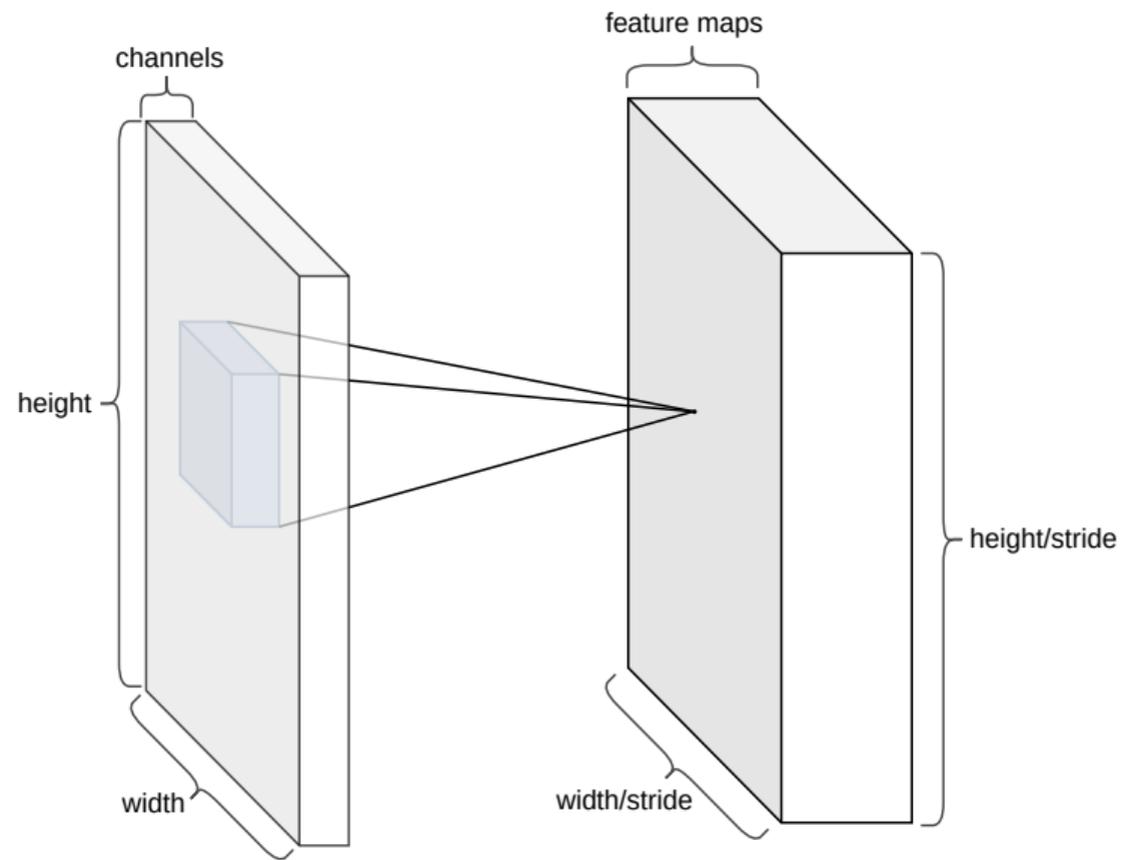
- To formulate the convolutions as matrix operations, the image pixels are duplicated and rearranged.





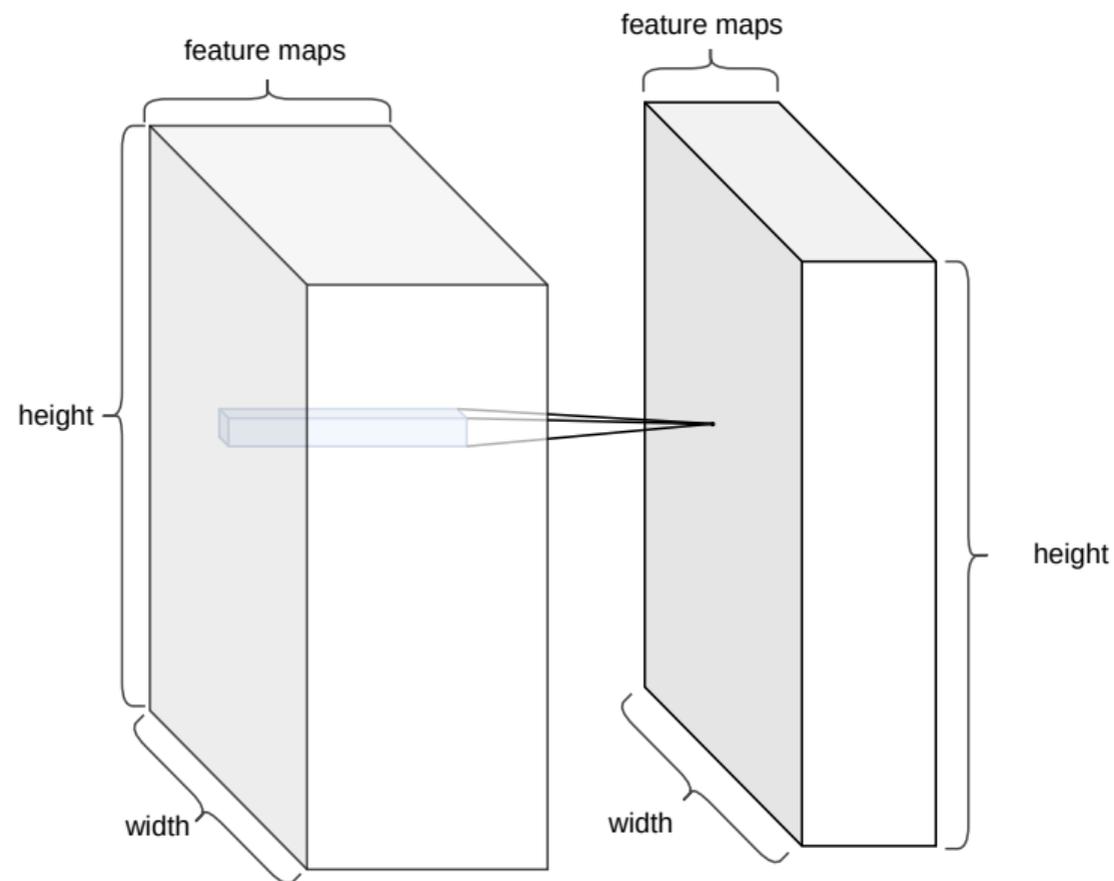
- Afterwards we multiply with a matrix that consists of the kernel values.
→ The computation of a forward/backward pass of convolutional layer becomes one big matrix operation.

Convolutional Layers



- The input to a convolutional layer is a tensor with $width \times height \times channels$
- The kernel is a four dimensional tensor with $nk \times ks \times ks \times c$, with number of kernels nk , the kernel size ks , and the number of channels c .
- The output is again a tensor $width' \times height' \times nk$, where the new width and height depend on padding and strides.
- The output channels are often referred to as feature maps.

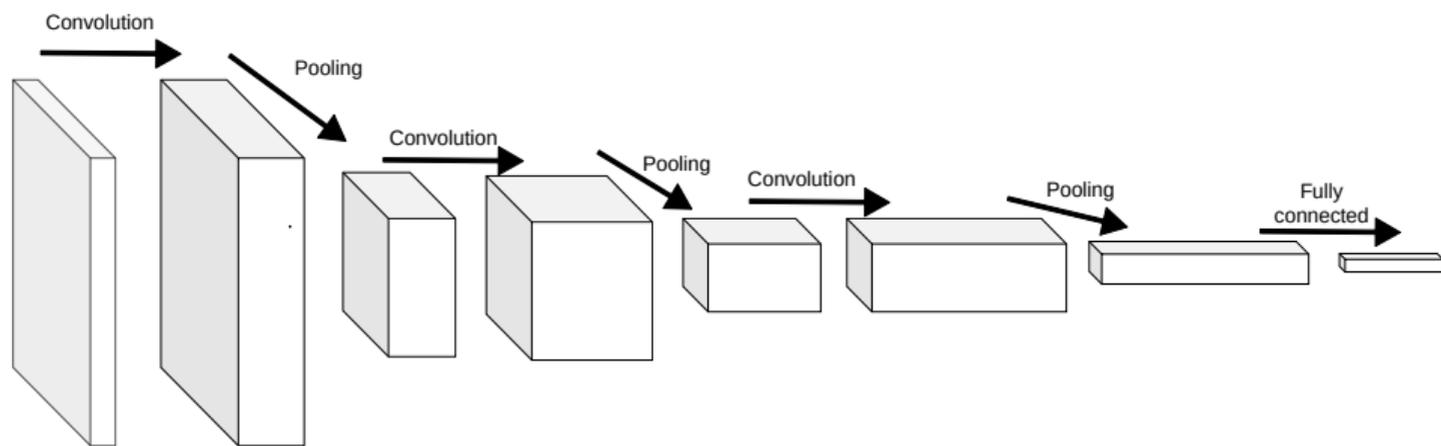
Convolutional Layers



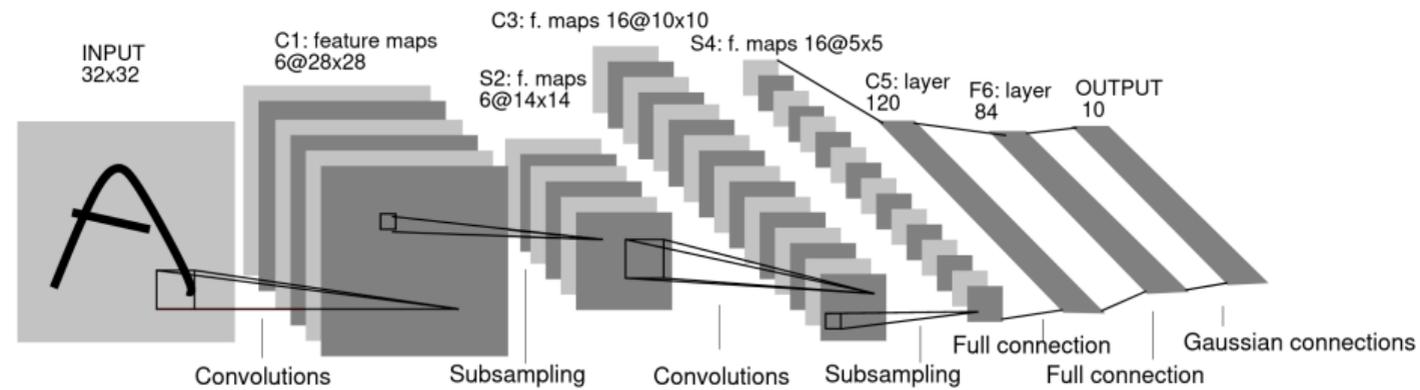
- Kernels with kernel size 1 can make sense, e.g. to reduce the number of feature maps.
- $fmaps' \times 1 \times 1 \times fmaps$ are called 1×1 convolutions.

Summary

-
-



- Gradient-based learning applied to document recognition, LeCun et al, 1998
- Classifies handwritten digits of the MNIST dataset.





- ImageNet is an image database organized according to the WordNet hierarchy (15 mio images).
- <https://www.image-net.org/>
- Widely used subset for ImageNet Large Scale Visual Recognition Challenge (ILSVRC): 1000 object classes, 1,281,167 training images, 50,000 validation images and 100,000 test images

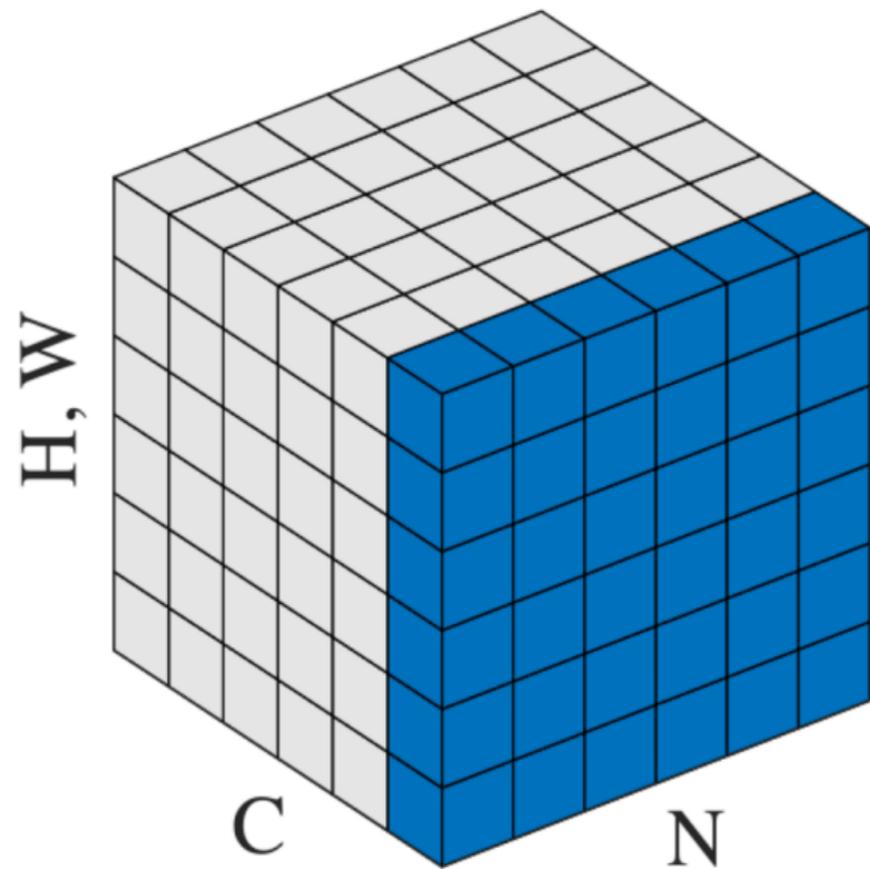
ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224×224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

Table 2: Number of parameters (in millions)

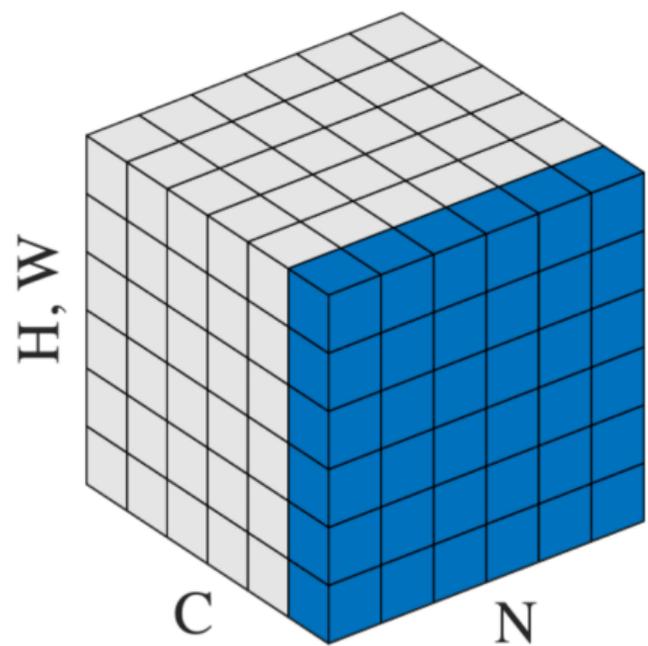
- Very Deep Convolutional Networks for Large-Scale Image Recognition, Simonyan & Zisserman, ICLR 2015
- Visual Geometry Group → VGG
- Depth matters, small kernels with size 3
- Still often used but really shouldn't.

-
-

- ▶ With deep networks and bounded activation functions gradients get very small.
- ▶ With unbounded activation functions gradients can explode.

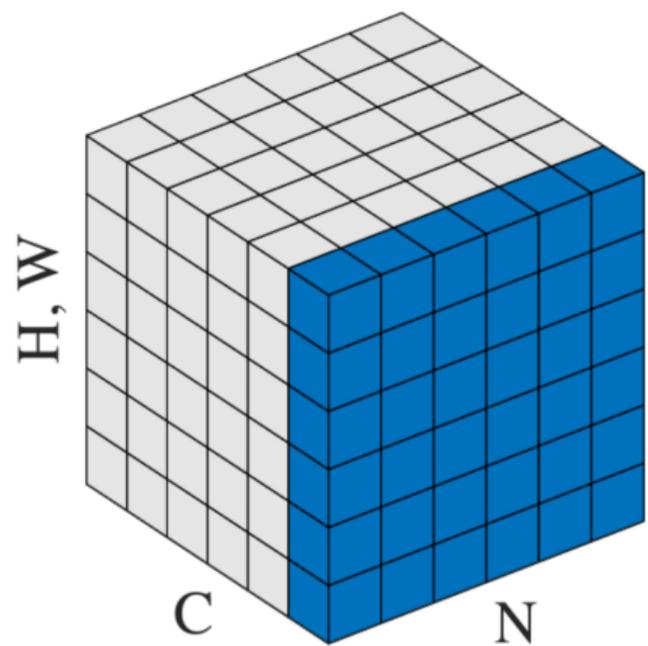


- Distribution of layer activations changes after every weight update! (Ioffe & Szegedy call this the internal covariate shift.)
- Lets normalize input to every layer!
- Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift, Ioffe & Szegedy, PLMR 2015
- Image from Group Normalization, Wu & He, Group normalization, 2018



$$x'_i = \frac{x_i - E[x_i]}{\sqrt{\text{Var}[x_i]}}$$

- It's as simple as the normalization of the input data. Almost ...
- What if mean and variance of activations matter?



$$x'_i = \frac{x_i - E[x_i]}{\sqrt{\text{Var}[x_i]}}$$
$$x'' = \gamma x' + \beta$$

- It's as simple as the normalization of the input data. Almost ...
- What if mean and variance of activations matter?
→ Add learnable parameters to modulate mean and variance!



- Usually inserted before the non-linearity

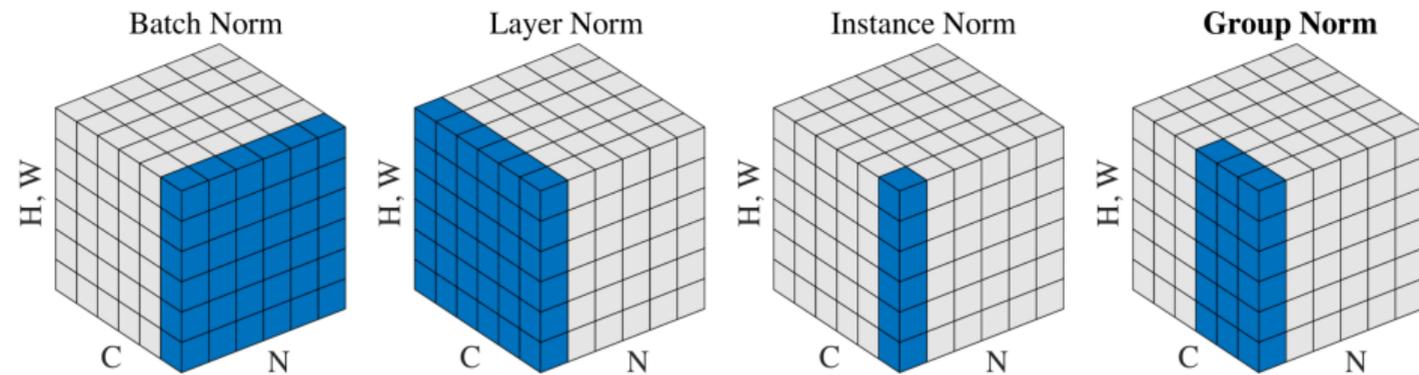
-
-

- ▶ Internal covariate shift is reduced
 - Training is more stable, higher learning rates possible
 - Contribution of samples in mini-batch to gradient harmonized
 - Input to non-linearity centered around zero
- ▶ Contribution to gradient of a sample depends on other samples in mini-batch
 - Regularization
 - In some cases detrimental

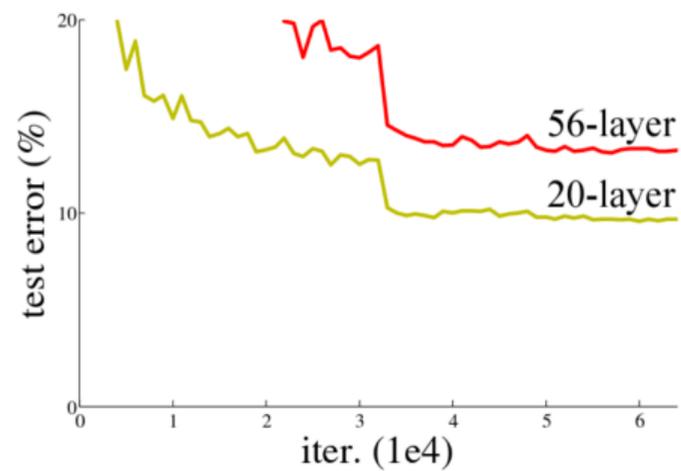
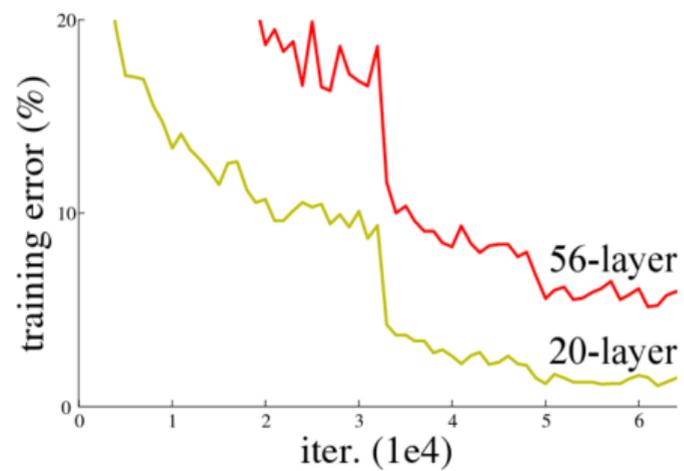
-
-

- ▶ Different behavior in training and test time
→ Often leads to bugs
- ▶ Adds a lot of complexity in recurrent networks
→ Every pass through a layer needs a dedicated batch-norm layer
- ▶ Depends on batch size (zero variance for single sample)

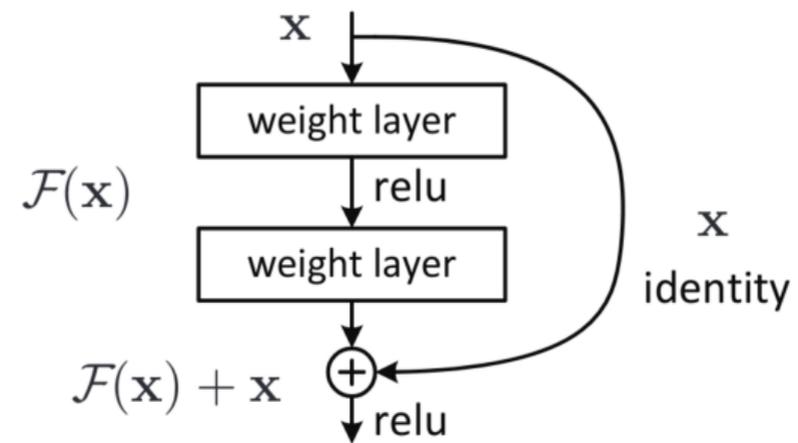
Alternative forms of normalization



- Layer Normalization, Ba et al, 2016
- Improved Texture Networks: Maximizing Quality and Diversity in Feed-forward Stylization and Texture Synthesis, Ulyanov et al, 2017
- Group Normalization, Wu & He, Group normalization, 2018
- Many more including combinations of these and weight normalization
- Image from Group Normalization, Wu & He, Group normalization, 2018



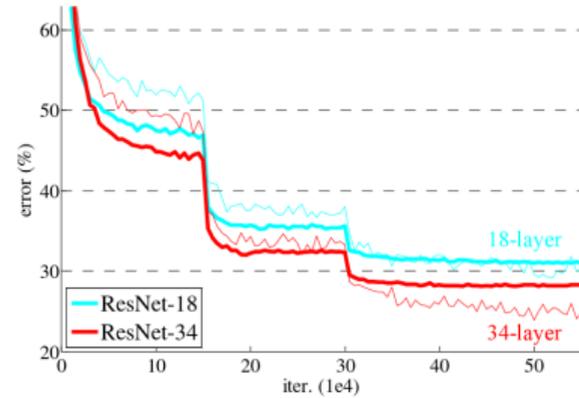
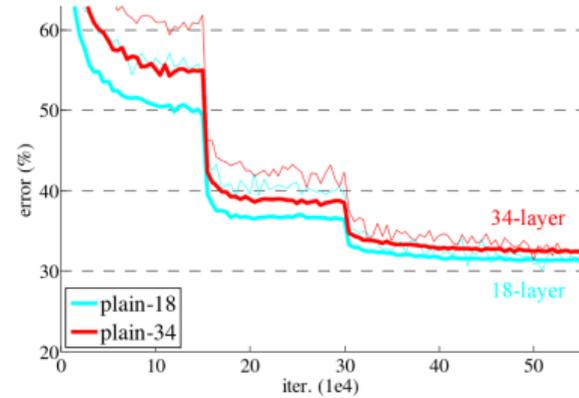
- VGG and others showed that accuracy increases with depth.
- Vanishing/exploding gradients are alleviated by normalization.
- But even with normalization, we can observe that training and test error start to increase again at a certain number of layers/depth.
- Deep Residual Learning for Image Recognition, He et al, CVPR 2016



- Idea: Shortcut layers, so learning the identity is setting weights to zero, which should be easier as actually learning the identity.

Residual Connections

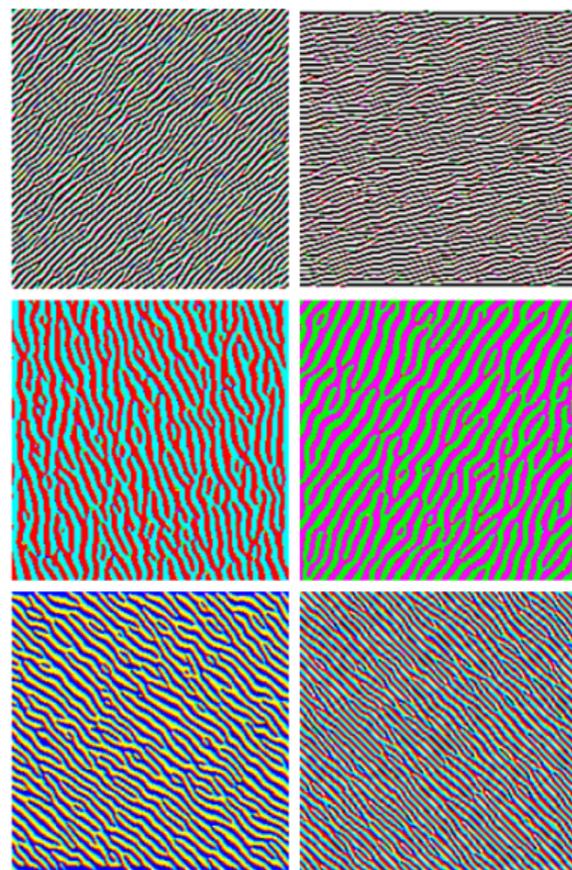
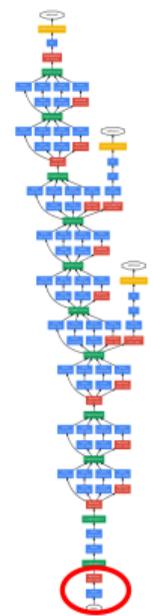
- Yay! Works even better!
- There is no limit to depth any more!
- Super human performance on ImageNet with a network with 152 layers.





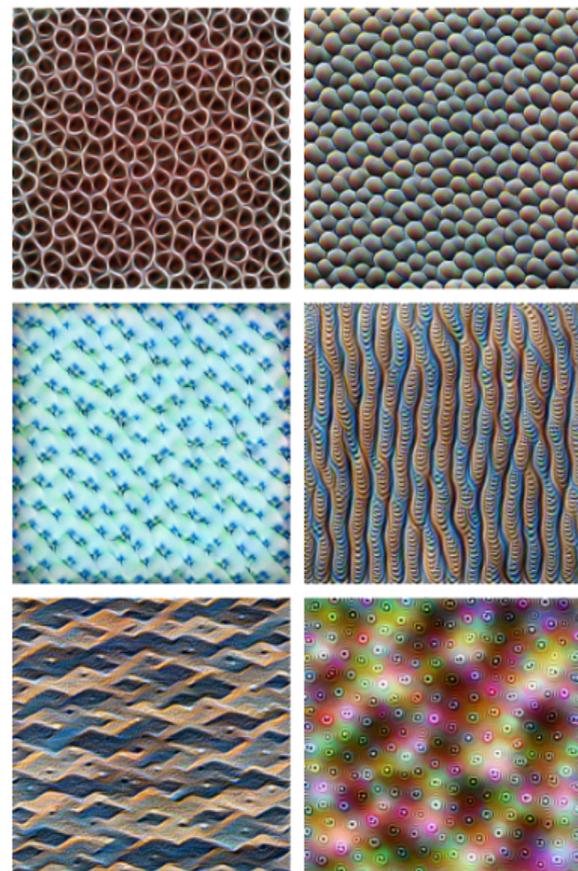
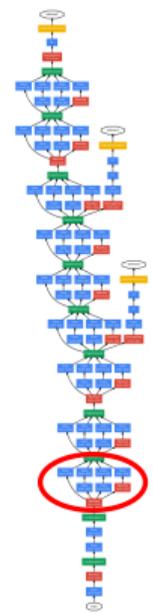
-
- ImageNet Classification with Deep Convolutional Neural Networks, Krizhevsky et al, NeurIPS 2012

Visualization: Early



- Going deeper with convolutions, Szegedy et al, CVPR 2015
- Feature Visualization, Olah et al, <https://distill.pub/2017/feature-visualization/>

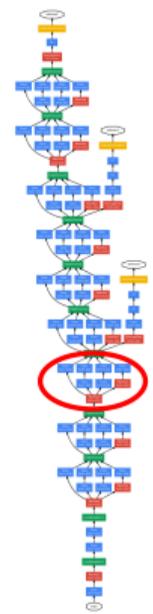
Visualization: Middle



■

■

Visualization: Middle

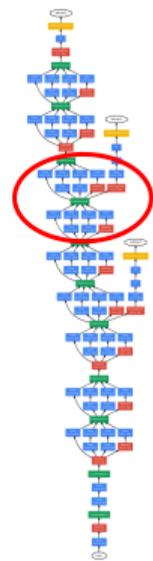


-
-

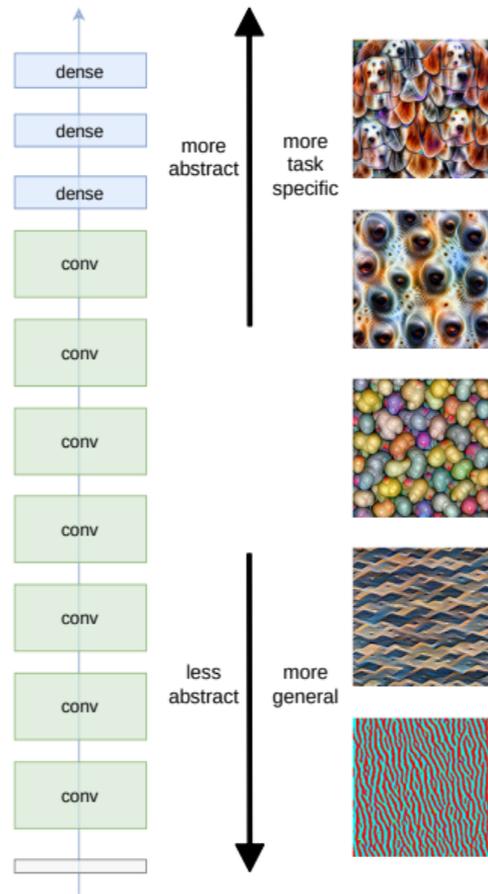
Visualization: Middle



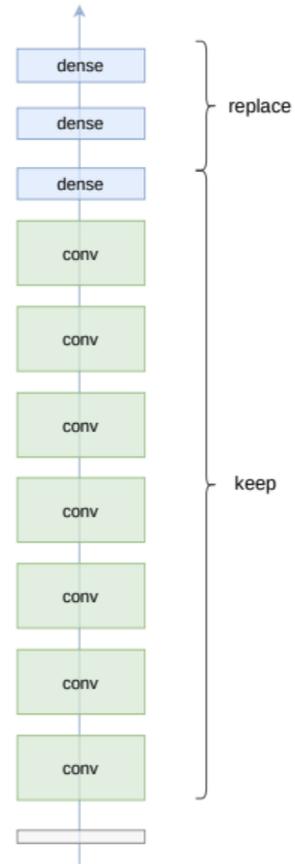
-
-



-
-

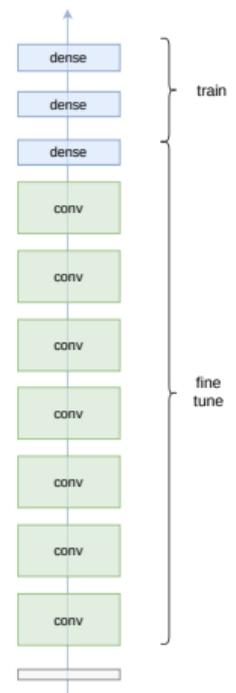


- Deep learning needs big data
- But what if we use the more general abilities a network learned for another task?
- Images from Feature Visualization, Olah et al, <https://distill.pub/2017/feature-visualization/>



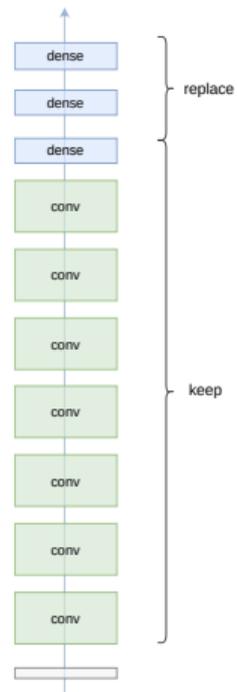
- We can transfer knowledge by replacing the specialized layers with randomly initialized layers and only train those!

Transfer Learning: fine tuning



- ▶ 1. Train the randomly initialized layers to convergence.
- ▶ 2. Unfreeze the some of the upper layers and continue training.

- Randomly initialized layers generate high gradients, which would destroy what was learned in the layers below.
- Fine tune with very small learning rate ($\approx .1 \times$ original lr)



- ▶ The more data we have for the target domain
 - the more layers we can replace.
 - the more layers we can fine tune.
- ▶ The higher the distance between original and target task
 - the more layers we may want to replace.
 - the more layers we need to fine tune.

- Give it a try, it works surprisingly well.
- Transfer learning has become the default initialization.
(In many frameworks, it's just an argument in a function call to initialize with weights trained on ImageNet.)
- Recent results show, that it is not always necessary.
(Rethinking ImageNet Pre-training, He et al, ICCV 2019)